



Death by a Thousand Leaks

What statically-analysing 370 Python
extensions looks like

Presented by
David Malcolm
<dmalcolm@redhat.com>

What is static analysis?

- Discovering properties of a program without running it
- Programs that analyze other programs
- Treating programs as data, rather than code
- In particular, automatically finding bugs in code

What kind of code will be
analyzed?

For this talk:

The C code of
Python extension modules

Prerequisites

- I'm going to assume basic familiarity with Python, and with either C or C++
- Hopefully you've used, debugged, or written a Python extension module in C (perhaps via SWIG or Cython)

Outline

- Intro to "cpychecker"
- How to run the tool on your own code
- How I ran the tool on lots of code
- What bugs came up frequently
- Recommendations on dealing with C and C++ from Python
- Q & A

cpychecker

**git clone **

git://git.fedorahosted.org/gcc-python-plugin.git

Docs: <http://tinyurl.com/cpychecker>

Part of my Python plugin for GCC

- 6500 lines of Python code implementing a static checker for C extension modules
- See also my PyCon US 2012 talk: Static analysis of Python extension modules using GCC
<https://us.pycon.org/2012/schedule/presentation/78/>

Reference counting



- For every object:
 - "what do I think my reference count is?" aka "ob_refcnt" (the object's view of how many pointers point to it) *versus*
 - the reality of how many pointers point to it
- As a C extension module author you must manually keep these in sync using Py_INCREF and Py_DECREF.

Reference counting



The two kinds of bugs:

- **ob_refcnt too high**
 - memory leaks (hence the title of this talk)
- **ob_refcnt too low**
 - **BOOM!!**

Checking reference counts

- For each path through the function and PyObject*, it determines:
 - what the reference count ought to be at the end of the function (based on how many pointers point to the object)
 - what the reference count *is*
- It will issues warnings for any that are incorrect.

File: **input.c**

Function: **test**

Error: **ob_refcnt of '*list' is 1 too high**

```
22 PyObject *
23 test(PyObject *self, PyObject *args)
24 {
25     PyObject *list;
26     PyObject *item;
27     list = PyList_New(1);
28     if (!list)
29         return NULL;
30     item = PyLong_FromLong(42);
31     /* This error handling is incorrect: it's missing an
32        invocation of Py_DECREF(list): */
33     if (!item)
34         return NULL;
35     /* This steals a reference to item; item is not leaked when we get here: */
36     PyList_SetItem(list, 0, item);
37     return list;
38 }
```

ob_refcnt of '*list' is 1 too high

was expecting final ob_refcnt to be $N + 0$ (for some unknown N)

but final ob_refcnt is $N + 1$

Limitations of the refcount checking

- purely intraprocedural
 - assumes every function returning a PyObject* returns a new reference, rather than a borrowed reference
 - (...although you can manually mark functions with non-standard behavior)
 - it knows about most of the CPython API and its rules

Limitations of the refcount checking (2)

- only tracks 0 and 1 times through any loop, to ensure that the analysis doesn't go on forever
- can be defeated by relatively simple code (turn up `--maxtrans` argument)

What it checks for (2)

It checks for the following along all of those code paths:

- Dereferencing a NULL pointer (e.g. using result of an allocator without checking the result is non-NULL)
- Passing NULL to CPython APIs that will crash on NULL

What it checks for (3)

- Usage of uninitialized local variables
- Dereferencing a pointer to freed memory
- Returning a pointer to freed memory
- Returning NULL without setting an exception

What it checks for (4)



It also does some simpler checking:

- type in calls to PyArg_ParseTuple et al
- types and NULL termination of PyMethodDef tables
- types and NULL termination of PyObject_Call{Function|Method}ObjArgs

What it doesn't check for

(patches welcome!)

- tp_traverse errors (which can mess up the garbage collector); missing it altogether, or omitting fields
- errors in GIL handling
- lock/release mismatches
- missed opportunities to release the GIL (e.g. compute-intensive functions; functions that wait on IO/syscalls)

What it can't check for



Does the code
"do the right thing"?

How to run it on your own code

```
git clone \
```

```
git://git.fedorahosted.org/gcc-python-plugin.git
```

Dependencies



(on Fedora)

```
sudo yum install \  
    gcc-plugin-devel \  
    python-devel \  
    python-six \  
    python-pygments \  
    graphviz
```

Building the checker



Building the checker:

make plugin

Checking that it's working:

make demo


```
demo.c: In function 'make_a_list_of_random_ints_badly':
demo.c:90:26: warning: Mismatching type in call to PyArg_ParseTuple with
argument 3 ("&count") had type
    "long int *" (pointing to 64 bits)
but was expecting
    "int *" (pointing to 32 bits)
for format code "i"
demo.c:102:1: warning: ob_refcnt of '*item' is 1 too high [enabled by de
demo.c:102:1: note: was expecting final ob_refcnt to be N + 1 (for some
demo.c:102:1: note: due to object being referenced by: PyListObject.ob_i
demo.c:102:1: note: but final ob_refcnt is N + 2
demo.c:97:14: note: PyLongObject allocated at:          item = PyLong_Fro
demo.c:90:26: note: when PyArg_ParseTuple() succeeds at:      if (!PyArg_
demo.c:90:8: note: taking False path at:      if (!PyArg_ParseTuple(args,
demo.c:94:10: note: reaching:      list = PyList_New(0);
demo.c:94:10: note: when PyList_New() succeeds at:      list = PyList_New
demo.c:96:5: note: when considering range: 1 <= count.0 <= 0x7fffffff at
demo.c:96:5: note: taking True path at:      for (i = 0; i < count; i++)
demo.c:97:31: note: reaching:          item = PyLong_FromLong(random());
demo.c:97:14: note: when PyLong_FromLong() succeeds at:      item = P
demo.c:97:14: note: ob_refcnt is now refs: 1 + N where N >= 0
demo.c:98:22: note: when PyList_Append() succeeds at:      PyList_App
demo.c:98:22: note: ob_refcnt is now refs: 2 + N where N >= 0
demo.c:98:22: note: '*item' is now referenced by 1 non-stack value(s): P
demo.c:96:5: note: when considering count.0 == (int)1 from demo.c:90 at:
demo.c:96:5: note: taking False path at:      for (i = 0; i < count; i++)
demo.c:101:5: note: reaching:      return list;
```

File: **input.c**

Function: **test**

Error: **ob_refcnt of '*list' is 1 too high**

```
22 PyObject *
23 test(PyObject *self, PyObject *args)
24 {
25     PyObject *list;
26     PyObject *item;
27     list = PyList_New(1);
28     if (!list)
29         return NULL;
30     item = PyLong_FromLong(42);
31     /* This error handling is incorrect: it's missing an
32        invocation of Py_DECREF(list): */
33     if (!item)
34         return NULL;
35     /* This steals a reference to item; item is not leaked when we get here: */
36     PyList_SetItem(list, 0, item);
37     return list;
38 }
```

ob_refcnt of '*list' is 1 too high

was expecting final ob_refcnt to be $N + 0$ (for some unknown N)

but final ob_refcnt is $N + 1$

Building with it

Distutils

```
CC=/path/to/built/plugin/gcc-with-cpychecker \  
python setup.py build
```

to set the environment variable

Makefiles

```
make CC=/path/to/built/plugin/gcc-with-cpychecker
```

to override the Makefile variable CC.

Let us know how you get on!

Mailing list:

- `gcc-python-plugin@lists.fedorahosted.org`
- See:
<https://fedorahosted.org/mailman/listinfo/gcc-python-plugin>

Analyze all the things!

- The goal: analyze all of the C Python extensions in a recent Linux distribution
 - Specifically: all of the Python 2 C code in Fedora 17
 - Every source rpm that builds something that links against libpython2.7
 - 370(ish) packages
- The reality:
 - Some unevenness in the data coverage, so take my numbers with a pinch of salt
 - Lots of bugfixing as I went...

Running cpychecker a lot

Scaling up to hundreds of projects:

- building via RPM
 - hides the distutils vs Makefile vs CMake etc
- "mock" builds
 - every build gets its own freshly-provisioned chroot
- Use this to reliably inject static analysis...

"mock-with-analysis"



Running checkers:

- cpychecker
- cppcheck
- clang-analyzer
- gcc warnings

<https://github.com/fedora-static-analysis/mock-with-analysis>

Scaling up (continued)

- separation of model from presentation
 - "Firehose" XML format:
 - <https://github.com/fedora-static-analysis/firehose>
- detect analyzers that fail or exceed 1 minute to run
- store the result in a database
- capture any sources mentioned in a report
- can also capture arbitrary data e.g. code metrics

```

382     free(expr); expr = NULL;
383 }
384 goto cleanup;
385
386 err:
387     error = errno;
388     py_decref(dict);
389     py_decref(tuple);
390     PyErr_SetString(PyExc_RuntimeError, strerror(error));
391 cleanup:
392     free(tmp);
393     free(rule_str);
394     free(expr);
395     errno = error;
396     return output;
397 }

```

ob_refcnt of return value is 1 too low

(emitted by cpychecker)

TODO: a detailed trace is available in the data model (not yet rendered in this report)

```

398
399 static int perform_ft_query(const apol_policy_t * policy, const options_t * opt, apol_vector_t
400 {
401     apol_filename_trans_query_t *ftq = NULL;
402     size_t i;
403     int error = 0;
404
405     if (!policy || !opt || !v) {
406         PyErr_SetString(PyExc_RuntimeError, strerror(EINVAL));
407         errno = EINVAL;

```



```
193 static PyObject *get_ipaddress(PyObject *self __unused, PyObject *args)
```

```
194 {
```

```
195     struct ifreq ifr;
```

```
196     int fd, err;
```

```
197     char *devname;
```

```
198     char ipaddr[20];
```

```
199
200     if (!PyArg_ParseTuple(args, "s", &devname))
```

Mismatching type in call to PyArg_ParseTuple with format code "s"

argument 3 ("&devname") had type "char *" but was expecting "const char *" for format code "s"
(emitted by cpychecker)

```
201         return NULL;
```

```
202
```

```
203     /* Setup our request structure. */
```

```
204     memset(&ifr, 0, sizeof(ifr));
```

```
205     strncpy(&ifr.ifr_name[0], devname, IFNAMSIZ);
```

```
206     ifr.ifr_name[IFNAMSIZ - 1] = 0;
```

```
207
```

```
193 static PyObject *get_ipaddress(PyObject *self __unused, PyObject *args)
```

```
194 {
```

```
195     struct ifreq ifr;
```

```
196     int fd, err;
```

```
197     const char *devname;
```

```
198     char ipaddr[20];
```

```
199
```

```
200     if (!PyArg_ParseTuple(args, "s", &devname))
```

```
201         return NULL;
```

```
202
```

```
203     /* Setup our request structure. */
```


```
204     memset(&ifr, 0, sizeof(ifr));
```

```
205     strncpy(&ifr.ifr_name[0], devname, IFNAMSIZ);
```

```
206     ifr.ifr_name[IFNAMSIZ - 1] = 0;
```

```
207
```


What are the 20 most commonly used Py/_Py entrypoints?

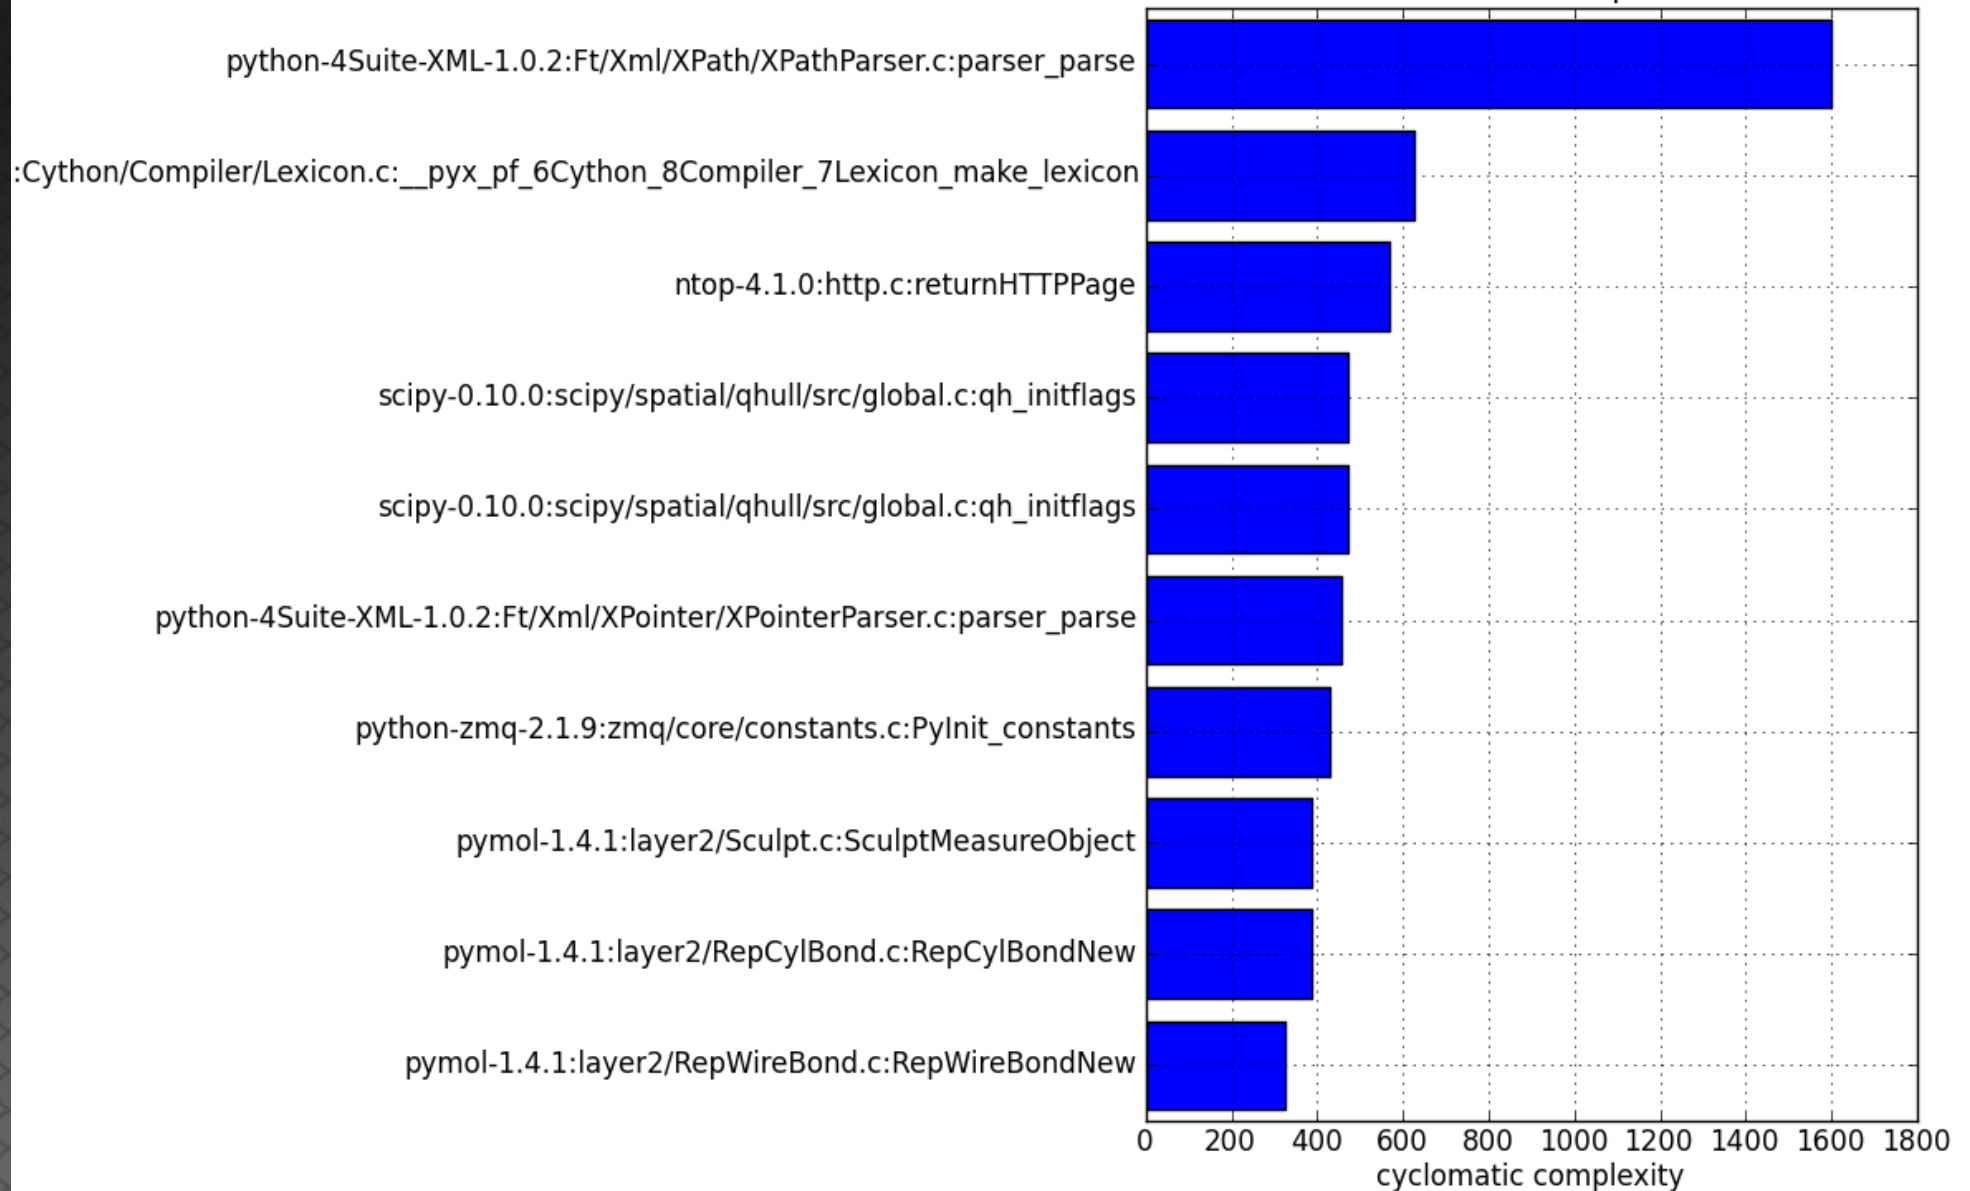


What are the least commonly used Py/_Py entrypoints?

- There are many with just 1 user, but most of these are false positives:
- about 50 actual CPython API entrypoints with just one user
- about 100 "entrypoints" due to other projects reusing the prefix


(see source code of this talk if you're interested in the data:
<https://github.com/davidmalcolm/PyCon-US-2013-Talk>

What were the 10 most complicated functions?

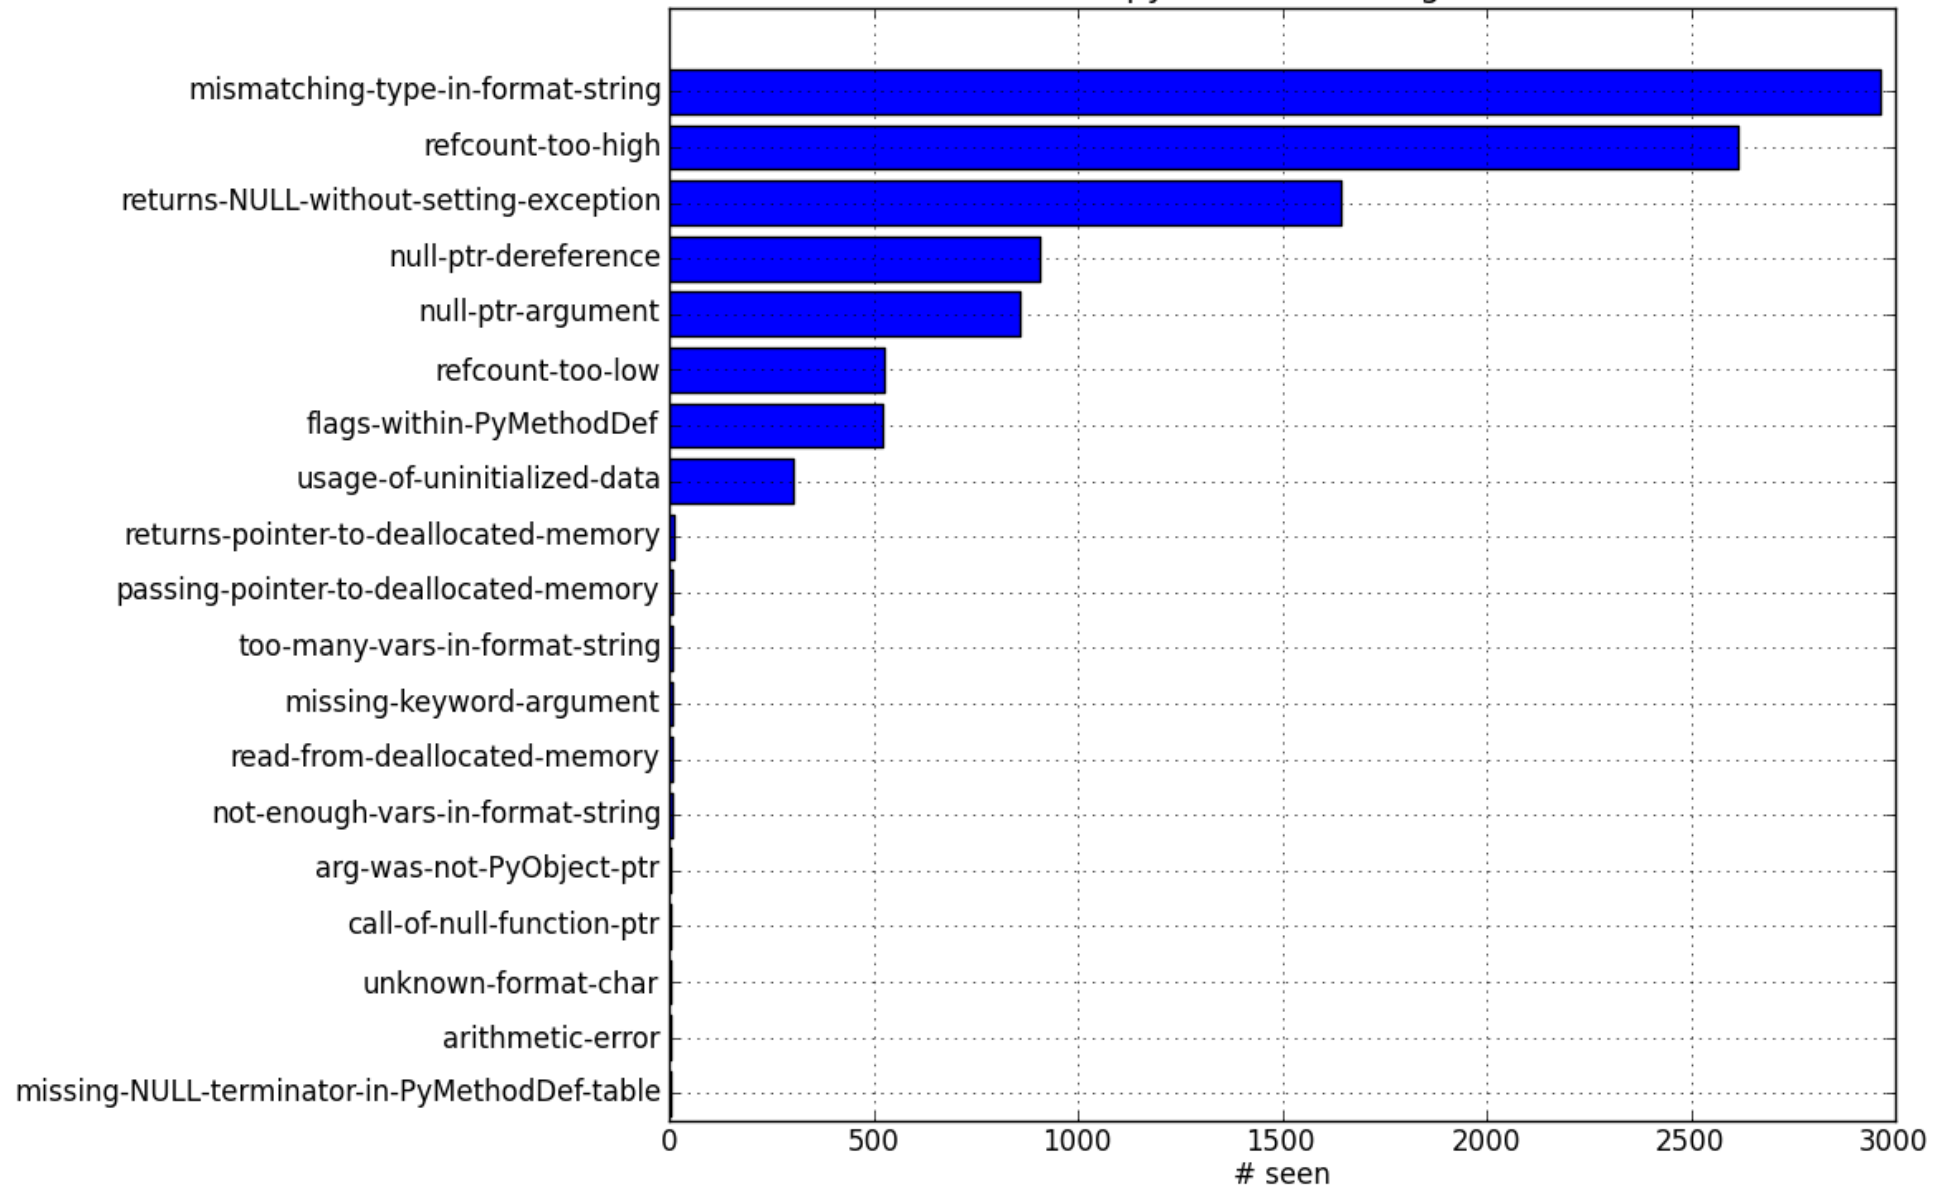


What did the analyzers
complain about?

Kinds of warning (top 20)



cpychecker warnings



Refcounting warnings



- refcount-too-high: 2614 times
- refcount-too-low: 524 times

Missing Py_INCREF() on Py_None

7% of the refcount-too-low warnings

(occurred 39 times (within 370 packages))

```
PyObject*  
some_method(PyObject *self, PyObject *args)  
{  
    [...snip...]  
  
    /* BUG: loses a reference to Py_None */  
    return Py_None;  
}
```

```
$ python script.py  
Fatal error: deallocating None
```

Fixing Py_INCREF on Py_None

```
PyObject*  
some_method(PyObject *self, PyObject *args)  
{  
    [...snip...]  
  
    /* Fixed version of the above: */  
    Py_RETURN_NONE;  
}
```

Reference leak in Py_BuildValue with "O"

```
/* BUG: reference leak: */  
return Py_BuildValue("Oi", some_object_we_own_a_ref_on, 42);
```

```
/* Fixed version of the above: */  
return Py_BuildValue("Ni", some_object_we_own_a_ref_on, 42);
```

```
/* If it's just one object, why use Py_BuildValue? */  
return some_object_we_own_a_ref_on;
```


1700+ places lacking error checking

- null-ptr-dereference: 907
- null-ptr-argument: 857

"goto" considered wonderful

```
{  
    PyObject *local0 = NULL;  
    PyObject *local1 = NULL;  
    PyObject *local2 = NULL;  
    /* etc */  
  
    local0 = PyFoo_DoBar();  
    if (!local0) goto error;  
  
    /* etc */  
  
    return result;  
  
error:  
    Py_XDECREF(local2);  
    Py_XDECREF(local1);  
    Py_XDECREF(local0);  
    return NULL;  
}
```

DO NOT DO THIS...

```
Py_XDECREF(PyObject_CallObject(callable, args));
```

How the compiler sees it...

```
do {
    if ((PyObject_CallObject(callable, args)) == ((void *)0))
        ;
    else
        do {
            if (--(PyObject_CallObject(callable, args)->ob_refcnt) != 0)
                ;
            else
                (*(PyObject_CallObject(callable, args)->ob_type)->tp_dealloc)
                PyObject_CallObject(callable, args);
        } while (0);
    } while (0);
```

Filed as <http://bugs.python.org/issue17206>

How the compiler sees it...

```
/* Call it once */
if ((PyObject_CallObject(callable, args)) != NULL) {
    /*
        If it doesn't raise an exception, leak the reference (BUG 1),
        and call it again (BUG 2).

        Assume that the second call doesn't raise an exception,
        otherwise segfault the interpreter (BUG 3),
        and DECREF the result, but don't deallocate if the refcount
        is zero (BUGS 4 and 5)
    */
    if (--(PyObject_CallObject(callable, args)->ob_refcnt) == 0) {
        /*
            If the refcount is zero, call it again! (BUG 6)
            Assume the result is non-NULL (otherwise segfaulting, BUG 7)
            and deallocate whatever you got back (even if the refcount
            is non-zero, BUG 8)
        */
        (*(PyObject_CallObject(callable, args)->ob_type)->tp_dealloc)
        /* and for good measure, call it agains (BUG 9)
           and leak a reference to the result (BUG 10) */
        PyObject_CallObject(callable, args);
    }
}
```

Filed as <http://bugs.python.org/issue17206>

The correct way to discard the result

```
PyObject *result;  
result = PyObject_CallObject(callable, args);  
Py_XDECREF(result);  
  
/* Presumably the caller will do something about any exception: */  
return (result != NULL) ? 0 : -1;
```

Dealing with C and C++ from Python

- Do you **really** need C?
- Can you get away with pure Python code?
- Consider using Cython
- ctypes is good, but has its own issues
- cffi?
- If you must use C, run cpychecker on your code

In conclusion

- Intro to "cpychecker"
- How to run the tool on your own code
- How I ran the tool on lots of code
- What bugs came up frequently
- Recommendations on dealing with C and C++ from Python

Thanks for listening!

Q & A

git clone \

[git://git.fedorahosted.org/gcc-python-plugin.git](https://git.fedorahosted.org/gcc-python-plugin.git)

- cpychecker's mailing list:
<https://fedorahosted.org/mailman/listinfo/gcc-python-plugin>
- This talk:
 - <https://github.com/davidmalcolm/PyCon-US-2013-Talk>